

# Kubernetes上安装KubeSphere

## 安装步骤

- ☐ 选择4核8G (master) 、8核16G (node1) 、8核16G (node2) 三台机器，按量付费进行实验，CentOS7.9
- ☐ 安装Docker
- ☐ 安装Kubernetes
- ☐ 安装KubeSphere前置环境
- ☐ 安装KubeSphere

## 1、安装Docker

```
sudo yum remove docker*
sudo yum install -y yum-utils

#配置docker的yum地址
sudo yum-config-manager \
--add-repo \
http://mirrors.aliyun.com/docker-ce/linux/centos/docker-ce.repo

#安装指定版本
sudo yum install -y docker-ce-20.10.7 docker-ce-cli-20.10.7 containerd.io-1.4.6

# 启动&开机启动docker
systemctl enable docker --now

# docker加速配置
sudo mkdir -p /etc/docker
sudo tee /etc/docker/daemon.json <<- 'EOF'
{
  "registry-mirrors": ["https://82m9ar63.mirror.aliyuncs.com"],
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
sudo systemctl daemon-reload
sudo systemctl restart docker
```

## 2、安装Kubernetes

### 1、基本环境

每个机器使用内网ip互通

每个机器配置自己的hostname，不能用localhost

```
#设置每个机器自己的hostname
hostnamectl set-hostname xxx

# 将 SELinux 设置为 permissive 模式（相当于将其禁用）
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config

#关闭swap
swapoff -a
sed -ri 's/.*swap.*/#&/' /etc/fstab

#允许 iptables 检查桥接流量
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
br_netfilter
EOF

cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
sudo sysctl --system
```

### 2、安装kubelet、kubeadm、kubectl

```
#配置k8s的yum源地址
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=http://mirrors.aliyun.com/kubernetes/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=0
repo_gpgcheck=0
gpgkey=http://mirrors.aliyun.com/kubernetes/yum/doc/yum-key.gpg
        http://mirrors.aliyun.com/kubernetes/yum/doc/rpm-package-key.gpg
EOF

#安装 kubelet, kubeadm, kubectl
sudo yum install -y kubelet-1.20.9 kubeadm-1.20.9 kubectl-1.20.9

#启动kubelet
sudo systemctl enable --now kubelet

#所有机器配置master域名
echo "172.31.0.4 k8s-master" >> /etc/hosts      #是master IP
```

## 3、初始化master节点

### 1、初始化

```
kubeadm init \  
--apiserver-advertise-address=172.31.0.4 \    #这个是master IP  
--control-plane-endpoint=k8s-master \  
--image-repository registry.cn-hangzhou.aliyuncs.com/lfy_k8s_images \  
--kubernetes-version v1.20.9 \  
--service-cidr=10.96.0.0/16 \  
--pod-network-cidr=192.168.0.0/16
```

### 2、记录关键信息

记录master执行完成后的日志

Your Kubernetes control-plane has initialized successfully!

To **start** using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, **if** you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "**kubectl apply -f [podnetwork].yaml**" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of control-plane nodes by copying certificate authorities

and **service** account keys on each **node** and **then** running the following as root:

```
kubeadm join k8s-master:6443 --token 3vckmv.1vr105xpyftbs177 \  
--discovery-token-ca-cert-hash  
sha256:1dc274fed24778f5c284229d9fcba44a5df11efba018f9664cf5e8ff77907240 \  
--control-plane
```

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join k8s-master:6443 --token 3vckmv.1vr105xpyftbs177 \  
--discovery-token-ca-cert-hash  
sha256:1dc274fed24778f5c284229d9fcba44a5df11efba018f9664cf5e8ff77907240
```

### 3、安装Calico网络插件

```
curl https://docs.projectcalico.org/manifests/calico.yaml -O  
  
kubectl apply -f calico.yaml
```

### 4、加入worker节点

---

## 3、安装KubeSphere前置环境

---

### 1、nfs文件系统

---

#### 1、安装nfs-server

```
# 在每个机器。  
yum install -y nfs-utils  
  
# 在master 执行以下命令  
echo "/nfs/data/ *(insecure,rw,sync,no_root_squash)" > /etc/exports
```

```
# 执行以下命令，启动 nfs 服务；创建共享目录  
mkdir -p /nfs/data
```

```
# 在master执行  
systemctl enable rpcbind  
systemctl enable nfs-server  
systemctl start rpcbind  
systemctl start nfs-server
```

```
# 使配置生效  
exportfs -r
```

```
#检查配置是否生效  
exportfs
```

## 2、配置nfs-client (选做)

```
#下面的命令在node1 node2 上执行
showmount -e 172.31.0.4      #master IP

mkdir -p /nfs/data

mount -t nfs 172.31.0.4:/nfs/data /nfs/data
```

## 3、配置默认存储

配置动态供应的默认存储类

```
## 创建了一个存储类
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: nfs-storage
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
provisioner: k8s-sigs.io/nfs-subdir-external-provisioner
parameters:
  archiveOnDelete: "true" ## 删除pv的时候，pv的内容是否要备份

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nfs-client-provisioner
  labels:
    app: nfs-client-provisioner
# replace with namespace where provisioner is deployed
namespace: default
spec:
  replicas: 1
  strategy:
    type: Recreate
  selector:
    matchLabels:
      app: nfs-client-provisioner
  template:
    metadata:
      labels:
        app: nfs-client-provisioner
    spec:
      serviceAccountName: nfs-client-provisioner
      containers:
        - name: nfs-client-provisioner
          image: registry.cn-hangzhou.aliyuncs.com/lfy_k8s_images/nfs-subdir-external-provisioner:v4.0.2
          # resources:
          #   limits:
          #     cpu: 10m
          #   requests:
```

```

    #      cpu: 10m
    volumeMounts:
      - name: nfs-client-root
        mountPath: /persistentvolumes
    env:
      - name: PROVISIONER_NAME
        value: k8s-sigs.io/nfs-subdir-external-provisioner
      - name: NFS_SERVER
        value: 172.31.0.4 ## 指定自己nfs服务器地址
      - name: NFS_PATH
        value: /nfs/data ## nfs服务器共享的目录
    volumes:
      - name: nfs-client-root
        nfs:
          server: 172.31.0.4 ##master节点
          path: /nfs/data
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: nfs-client-provisioner
  # replace with namespace where provisioner is deployed
  namespace: default
---
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: nfs-client-provisioner-runner
rules:
  - apiGroups: [""]
    resources: ["nodes"]
    verbs: ["get", "list", "watch"]
  - apiGroups: [""]
    resources: ["persistentvolumes"]
    verbs: ["get", "list", "watch", "create", "delete"]
  - apiGroups: [""]
    resources: ["persistentvolumeclaims"]
    verbs: ["get", "list", "watch", "update"]
  - apiGroups: ["storage.k8s.io"]
    resources: ["storageclasses"]
    verbs: ["get", "list", "watch"]
  - apiGroups: [""]
    resources: ["events"]
    verbs: ["create", "update", "patch"]
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: run-nfs-client-provisioner
subjects:
  - kind: ServiceAccount
    name: nfs-client-provisioner
    # replace with namespace where provisioner is deployed
    namespace: default
roleRef:
  kind: ClusterRole
  name: nfs-client-provisioner-runner
  apiGroup: rbac.authorization.k8s.io

```

```

---
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: leader-locking-nfs-client-provisioner
  # replace with namespace where provisioner is deployed
  namespace: default
rules:
- apiGroups: [""]
  resources: ["endpoints"]
  verbs: ["get", "list", "watch", "create", "update", "patch"]
---
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: leader-locking-nfs-client-provisioner
  # replace with namespace where provisioner is deployed
  namespace: default
subjects:
- kind: ServiceAccount
  name: nfs-client-provisioner
  # replace with namespace where provisioner is deployed
  namespace: default
roleRef:
  kind: Role
  name: leader-locking-nfs-client-provisioner
  apiGroup: rbac.authorization.k8s.io

```

```

#确认配置是否生效
kubectl get sc

```

## 2、metrics-server

集群指标监控组件

```

apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    k8s-app: metrics-server
    rbac.authorization.k8s.io/aggregate-to-admin: "true"
    rbac.authorization.k8s.io/aggregate-to-edit: "true"
    rbac.authorization.k8s.io/aggregate-to-view: "true"
  name: system:aggregated-metrics-reader
rules:
- apiGroups:

```

```

- metrics.k8s.io
resources:
- pods
- nodes
verbs:
- get
- list
- watch
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    k8s-app: metrics-server
  name: system:metrics-server
rules:
- apiGroups:
  - ""
  resources:
  - pods
  - nodes
  - nodes/stats
  - namespaces
  - configmaps
  verbs:
  - get
  - list
  - watch
---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server-auth-reader
  namespace: kube-system
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: extension-apiserver-authentication-reader
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server:system:auth-delegator
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:auth-delegator
subjects:
- kind: ServiceAccount
  name: metrics-server

```



```

    namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    k8s-app: metrics-server
  name: system:metrics-server
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:metrics-server
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system
---
apiVersion: v1
kind: Service
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: kube-system
spec:
  ports:
  - name: https
    port: 443
    protocol: TCP
    targetPort: https
  selector:
    k8s-app: metrics-server
---
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: kube-system
spec:
  selector:
    matchLabels:
      k8s-app: metrics-server
  strategy:
    rollingUpdate:
      maxUnavailable: 0
  template:
    metadata:
      labels:
        k8s-app: metrics-server
    spec:
      containers:
      - args:
        - --cert-dir=/tmp
        - --kubelet-insecure-tls
        - --secure-port=4443
        - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname

```

```

- --kubelet-use-node-status-port
image: registry.cn-hangzhou.aliyuncs.com/lfy_k8s_images/metrics-
server:v0.4.3
imagePullPolicy: IfNotPresent
livenessProbe:
  failureThreshold: 3
  httpGet:
    path: /livez
    port: https
    scheme: HTTPS
    periodSeconds: 10
name: metrics-server
ports:
- containerPort: 4443
  name: https
  protocol: TCP
readinessProbe:
  failureThreshold: 3
  httpGet:
    path: /readyz
    port: https
    scheme: HTTPS
    periodSeconds: 10
securityContext:
  readOnlyRootFilesystem: true
  runAsNonRoot: true
  runAsUser: 1000
volumeMounts:
- mountPath: /tmp
  name: tmp-dir
nodeSelector:
  kubernetes.io/os: linux
priorityClassName: system-cluster-critical
serviceAccountName: metrics-server
volumes:
- emptyDir: {}
  name: tmp-dir
---
apiVersion: apiregistration.k8s.io/v1
kind: APIService
metadata:
  labels:
    k8s-app: metrics-server
  name: v1beta1.metrics.k8s.io
spec:
  group: metrics.k8s.io
  groupPriorityMinimum: 100
  insecureSkipTLSVerify: true
  service:
    name: metrics-server
    namespace: kube-system
  version: v1beta1
  versionPriority: 100

```

## 4、安装KubeSphere

<https://kubesphere.com.cn/>

### 1、下载核心文件

如果下载不到，请复制附录的内容

```
wget https://github.com/kubesphere/ks-installer/releases/download/v3.1.1/kubesphere-installer.yaml
```

```
wget https://github.com/kubesphere/ks-installer/releases/download/v3.1.1/cluster-configuration.yaml
```

### 2、修改cluster-configuration

在 cluster-configuration.yaml中指定我们需要开启的功能

参照官网“启用可插拔组件”

<https://kubesphere.com.cn/docs/pluggable-components/overview/>

### 3、执行安装

```
kubectl apply -f kubesphere-installer.yaml
```

```
kubectl apply -f cluster-configuration.yaml
```

### 4、查看安装进度

```
kubectl logs -n kubesphere-system $(kubectl get pod -n kubesphere-system -l app=ks-install -o jsonpath='{.items[0].metadata.name}') -f
```

访问任意机器的 30880端口

账号： admin

密码： P@88w0rd

解决etcd监控证书找不到问题

```
kubectl -n kubesphere-monitoring-system create secret generic kube-etcd-client-certs --from-file=etcd-client-ca.crt=/etc/kubernetes/pki/etcd/ca.crt --from-file=etcd-client.crt=/etc/kubernetes/pki/apiserver-etcd-client.crt --from-file=etcd-client.key=/etc/kubernetes/pki/apiserver-etcd-client.key
```

# 附录

## 1、kubesphere-installer.yaml

```
---
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: clusterconfigurations.installer.kubesphere.io
spec:
  group: installer.kubesphere.io
  versions:
    - name: v1alpha1
      served: true
      storage: true
  scope: Namespaced
  names:
    plural: clusterconfigurations
    singular: clusterconfiguration
    kind: ClusterConfiguration
    shortNames:
      - cc

---
apiVersion: v1
kind: Namespace
metadata:
  name: kubesphere-system

---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: ks-installer
  namespace: kubesphere-system

---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: ks-installer
rules:
- apiGroups:
  - ""
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - apps
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - extensions
  resources:
```

```
- '*'
verbs:
- '*'
- apiGroups:
  - batch
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - rbac.authorization.k8s.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - apiregistration.k8s.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - apiextensions.k8s.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - tenant.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - certificates.k8s.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - devops.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - monitoring.coreos.com
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - logging.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
```

```
- jaegertracing.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- storage.k8s.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- admissionregistration.k8s.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- policy
resources:
- '*'

verbs:
- '*'

- apiGroups:
- autoscaling
resources:
- '*'

verbs:
- '*'

- apiGroups:
- networking.istio.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- config.istio.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- iam.kubesphere.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- notification.kubesphere.io
resources:
- '*'

verbs:
- '*'

- apiGroups:
- auditing.kubesphere.io
resources:
- '*'

verbs:
```

```
- '*'
- apiGroups:
  - events.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - core.kubefed.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - installer.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - storage.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - security.istio.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - monitoring.kiali.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - kiali.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - networking.k8s.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - kubeedge.kubesphere.io
  resources:
  - '*'
  verbs:
  - '*'
- apiGroups:
  - types.kubefed.io
  resources:
```

```

- '*'
verbs:
- '*'

---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: ks-installer
subjects:
- kind: ServiceAccount
  name: ks-installer
  namespace: kubescape-system
roleRef:
  kind: ClusterRole
  name: ks-installer
  apiGroup: rbac.authorization.k8s.io

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ks-installer
  namespace: kubescape-system
  labels:
    app: ks-install
spec:
  replicas: 1
  selector:
    matchLabels:
      app: ks-install
  template:
    metadata:
      labels:
        app: ks-install
    spec:
      serviceAccountName: ks-installer
      containers:
        - name: installer
          image: kubescape/ks-installer:v3.1.1
          imagePullPolicy: "Always"
          resources:
            limits:
              cpu: "1"
              memory: 1Gi
            requests:
              cpu: 20m
              memory: 100Mi
          volumeMounts:
            - mountPath: /etc/localtime
              name: host-time
      volumes:
        - hostPath:
            path: /etc/localtime
            type: ""
            name: host-time

```



## 2、cluster-configuration.yaml

```
---
apiVersion: installer.kubesphere.io/v1alpha1
kind: ClusterConfiguration
metadata:
  name: ks-installer
  namespace: kubesphere-system
  labels:
    version: v3.1.1
spec:
  persistence:
    storageClass: ""      # If there is no default StorageClass in your
cluster, you need to specify an existing StorageClass here.
  authentication:
    jwtSecret: ""        # Keep the jwtSecret consistent with the Host
Cluster. Retrieve the jwtSecret by executing "kubectl -n kubesphere-system get cm
kubesphere-config -o yaml | grep -v "apiVersion" | grep jwtSecret" on the Host
Cluster.
    local_registry: ""   # Add your private registry address if it is needed.
  etcd:
    monitoring: true     # Enable or disable etcd monitoring dashboard
installation. You have to create a Secret for etcd before you enable it.
    endpointIps: 172.31.0.4 # etcd cluster EndpointIps. It can be a bunch of
IPs here.
    port: 2379           # etcd port.
    tlsEnable: true
  common:
    redis:
      enabled: true
    openldap:
      enabled: true
    minioVolumeSize: 20Gi # Minio PVC size.
    openldapVolumeSize: 2Gi # openldap PVC size.
    redisVolumeSize: 2Gi # Redis PVC size.
    monitoring:
      # type: external # Whether to specify the external prometheus stack, and
need to modify the endpoint at the next line.
      endpoint: http://prometheus-operated.kubesphere-monitoring-system.svc:9090
# Prometheus endpoint to get metrics data.
      es: # Storage backend for logging, events and auditing.
        # elasticsearchMasterReplicas: 1 # The total number of master nodes.
Even numbers are not allowed.
        # elasticsearchDataReplicas: 1 # The total number of data nodes.
        elasticsearchMasterVolumeSize: 4Gi # The volume size of Elasticsearch
master nodes.
        elasticsearchDataVolumeSize: 20Gi # The volume size of Elasticsearch
data nodes.
        logMaxAge: 7 # Log retention time in built-in
Elasticsearch. It is 7 days by default.
        elkPrefix: logstash # The string making up index names. The
index name will be formatted as ks-<elk_prefix>-log.
    basicAuth:
      enabled: false
      username: ""
```

```

    password: ""
    externalElasticsearchUrl: ""
    externalElasticsearchPort: ""
console:
    enableMultiLogin: true # Enable or disable simultaneous logins. It allows
different users to log in with the same account at the same time.
    port: 30880
alerting: # (CPU: 0.1 Core, Memory: 100 MiB) It enables users
to customize alerting policies to send messages to receivers in time with
different time intervals and alerting levels to choose from.
    enabled: true # Enable or disable the KubeSphere Alerting System.
    # thanosruler:
    #   replicas: 1
    #   resources: {}
auditing: # Provide a security-relevant chronological set of
records, recording the sequence of activities happening on the platform,
initiated by different tenants.
    enabled: true # Enable or disable the KubeSphere Auditing Log
System.
devops: # (CPU: 0.47 Core, Memory: 8.6 G) Provide an out-of-
the-box CI/CD system based on Jenkins, and automated workflow tools including
Source-to-Image & Binary-to-Image.
    enabled: true # Enable or disable the KubeSphere DevOps System.
    jenkinsMemoryLim: 2Gi # Jenkins memory limit.
    jenkinsMemoryReq: 1500Mi # Jenkins memory request.
    jenkinsVolumeSize: 8Gi # Jenkins volume size.
    jenkinsJavaOpts_Xms: 512m # The following three fields are JVM parameters.
    jenkinsJavaOpts_Xmx: 512m
    jenkinsJavaOpts_MaxRAM: 2g
events: # Provide a graphical web console for Kubernetes
Events exporting, filtering and alerting in multi-tenant Kubernetes clusters.
    enabled: true # Enable or disable the KubeSphere Events System.
    ruler:
        enabled: true
        replicas: 2
logging: # (CPU: 57 m, Memory: 2.76 G) Flexible logging
functions are provided for log query, collection and management in a unified
console. Additional log collectors can be added, such as Elasticsearch, Kafka and
Fluentd.
    enabled: true # Enable or disable the KubeSphere Logging System.
    logsidecar:
        enabled: true
        replicas: 2
metrics_server: # (CPU: 56 m, Memory: 44.35 MiB) It enables
HPA (Horizontal Pod Autoscaler).
    enabled: false # Enable or disable metrics-server.
monitoring:
    storageClass: "" # If there is an independent StorageClass
you need for Prometheus, you can specify it here. The default StorageClass is
used by default.
    # prometheusReplicas: 1 # Prometheus replicas are responsible for
monitoring different segments of data source and providing high availability.
    prometheusMemoryRequest: 400Mi # Prometheus request memory.
    prometheusVolumeSize: 20Gi # Prometheus PVC size.
    # alertmanagerReplicas: 1 # AlertManager Replicas.
multicluster:
    clusterRole: none # host | member | none # You can install a solo cluster,
or specify it as the Host or Member Cluster.

```

```

network:
  networkpolicy: # Network policies allow network isolation within the same
cluster, which means firewalls can be set up between certain instances (Pods).
    # Make sure that the CNI network plugin used by the cluster supports
NetworkPolicy. There are a number of CNI network plugins that support
NetworkPolicy, including Calico, Cilium, Kube-router, Romana and Weave Net.
    enabled: true # Enable or disable network policies.
  ippool: # Use Pod IP Pools to manage the Pod network address space. Pods to
be created can be assigned IP addresses from a Pod IP Pool.
    type: calico # Specify "calico" for this field if Calico is used as your
CNI plugin. "none" means that Pod IP Pools are disabled.
  topology: # Use Service Topology to view Service-to-Service communication
based on Weave Scope.
    type: none # Specify "weave-scope" for this field to enable Service
Topology. "none" means that Service Topology is disabled.
  openpitrix: # An App Store that is accessible to all platform tenants. You can
use it to manage apps across their entire lifecycle.
  store:
    enabled: true # Enable or disable the KubeSphere App Store.
  servicemesh:
    # (0.3 Core, 300 MiB) Provide fine-grained traffic
management, observability and tracing, and visualized traffic topology.
    enabled: true # Base component (pilot). Enable or disable KubeSphere
Service Mesh (Istio-based).
  kubeedge:
    # Add edge nodes to your cluster and deploy workloads on
edge nodes.
    enabled: true # Enable or disable KubeEdge.
  cloudcore:
    nodeSelector: {"node-role.kubernetes.io/worker": ""}
    tolerations: []
    cloudhubPort: "10000"
    cloudhubQuicPort: "10001"
    cloudhubHttpsPort: "10002"
    cloudstreamPort: "10003"
    tunnelPort: "10004"
    cloudHub:
      advertiseAddress: # At least a public IP address or an IP address which
can be accessed by edge nodes must be provided.
        - "" # Note that once KubeEdge is enabled, CloudCore will
malfunction if the address is not provided.
      nodeLimit: "100"
    service:
      cloudhubNodePort: "30000"
      cloudhubQuicNodePort: "30001"
      cloudhubHttpsNodePort: "30002"
      cloudstreamNodePort: "30003"
      tunnelNodePort: "30004"
  edgewatcher:
    nodeSelector: {"node-role.kubernetes.io/worker": ""}
    tolerations: []
  edgewatcherAgent:
    nodeSelector: {"node-role.kubernetes.io/worker": ""}
    tolerations: []

```

